

```
In [6]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Load dataset
df = pd.read_csv('../Datasets/diabetes.csv')
M=len(df)
print(f"Loaded {M} samples.") # rather than a random "768" I will add additional inf
# from here I can go ahead and make the data provided easier to read and make sure
print(df.head()) # from the df.head() I went ahead and added print to make it show
```

Loaded 768 samples.

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
In [2]: df.describe()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.481517	33.421381	0.516742
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.332641	11.952609	0.501390
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	19.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243000	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.369000	31.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.625000	36.000000	0.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.425000	81.000000	1.000000

```
In [12]: # Suppress convergence warnings
import warnings
from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore", category=ConvergenceWarning)
```

```

# Prepare data
X = df.drop("Outcome", axis=1)
y = df["Outcome"]

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split into train and test sets (80/20)
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)

# Initialize model
model = LogisticRegression(max_iter=1000)

# Track loss and accuracy during training
losses, accuracies = [], []

for i in range(1, 101):
    model.max_iter = i
    model.fit(X_train, y_train)

    logits = model.decision_function(X_train)
    loss = np.mean(np.log(1 + np.exp(-y_train * logits)))
    acc = accuracy_score(y_train, model.predict(X_train))

    losses.append(loss)
    accuracies.append(acc)

# Evaluate final model on test set
y_pred = model.predict(X_test)
acc_test = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

```

```

In [13]: # Plot training metrics
plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(losses, label='Loss')
plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.title('Training Loss')
plt.legend()


# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(accuracies, label='Accuracy', color='green')
plt.xlabel('Iteration')
plt.ylabel('Accuracy')
plt.title('Training Accuracy')
plt.legend()


```

```
plt.tight_layout()
plt.show()

# Confusion matrix
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blues',
            xticklabels=["No Diabetes", "Diabetes"],
            yticklabels=["No Diabetes", "Diabetes"]))
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# Print evaluation metrics
print(f"Test Accuracy : {acc_test:.4f}")
print(f"Precision      : {prec:.4f}")
print(f"Recall          : {rec:.4f}")
print(f"F1 Score       : {f1:.4f}")
```





Test Accuracy : 0.7143  
Precision : 0.6087  
Recall : 0.5185  
F1 Score : 0.5600

```
In [10]: # I found this plot function and it is really crazy how detailed it gets  
sns.pairplot(df, hue = 'Outcome')
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x14e2e640770>
```



In [ ]: